

Attention Is All You Need

Abstract

We propose the Transformer architecture for sequence transduction.

1 Introduction

Recurrent models have been dominant in sequence modeling.

3 Methods

The Transformer uses multi-head self-attention.